

Alpinist

an Annotation-Aware GPU Program Optimizer

Ömer Şakar¹ Mohsen Safari¹ Marieke Huisman¹ Anton Wijs²

Formal Methods and Tools, University of Twente, Enschede, The Netherlands


Software Engineering & Technology, Eindhoven University of Technology, Eindhoven, The Netherlands




April 7, 2022




GPU Program Development Cycle




GPU Program Development Cycle




GPU Program Verification



GPU Program Verification



ALPINIST an Annotation-Aware GPU Program Optimizer



Abstract GPU Program

```
1 void simple_example(int[] a, int N) {
2     par kernel1 (int tid = 0 .. a.length) {
3         a[tid] = 0;
4     }
5
6     par kernel2(int tid = 0 .. a.length) {
7         for(int k = 0; k < N; k++)
8         {
9             if (tid != a.length-1) {
10                 a[tid+1] = a[tid+1] + tid;
11             } else {
12                 a[0] = a[0] + tid;
13             }
14         }
15     }
16 }
```

Abstract GPU Program

```
1 void simple_example(int[] a,int N) {
2     par fused_kernels(int fused_tid = 0 .. a.length) {
3         a [ fused_tid ] = 0;
4
5         barrier(fused_kernels)
6
7         for(int k = 0;k < N;k ++){
8             if (fused_tid != a.length - 1) {
9                 a [ (fused_tid + 1) ] = a [ (fused_tid + 1) ] + fused_tid;
10            } else {
11                a [ 0 ] = a [ 0 ] + fused_tid;
12            }
13        }
14    }
15 }
```

Abstract GPU Program with Annotations

```
1 ...
2 void simple_example(int[] a, int N) {
3     par kernel1 (int tid = 0 .. a.length)
4         /*@ context Perm(a[tid], 1);
5          ensures a[tid] == 0; @*/
6         { a[tid] = 0; }
7
8     ...
9 } } }
```

Abstract GPU Program with Annotations

```
1 ...
2 void simple_example(int[] a, int N) {
3     par kernel1 (int tid = 0 .. a.length)
4         /*@ context Perm(a[tid], 1);
5          ensures a[tid] == 0; @*/
6         { a[tid] = 0; }
7
8     ...
9 } } }
```

Abstract GPU Program with Annotations


```
1 ...
2 void simple_example(int[] a, int N) {
3     par kernel1 (int tid = 0 .. a.length)
4         /*@ context Perm(a[tid], 1\2);
5            ensures a[tid] == 0; @*/
6         { a[tid] = 0; }
7 ...
8 ...
9 } } }
```

Abstract GPU Program with Annotations

```
1 ...
2 void simple_example(int[] a, int N) {
3     par kernel1 (int tid = 0 .. a.length)
4     /*@ context Perm(a[tid], 1);
5      ensures a[tid] == 0;  @*/
6     { a[tid] = 0; }
7
8 ...
9 } } }
```

```
1   ...
2   void simple_example(int[] a, int N) {
3       par kernel2 (int tid = 0 .. a.length)
4           /*@ context Perm(a[tid], 1);
5            ensures a[tid] == 0; */
6           { a[tid] = 0; }
7
8       par kernel2 (int tid = 0 .. a.length)
9           /*@ context tid != a.length-1 ? Perm(a[tid+1], 1) : Perm(a[0], 1);
10          requires tid != a.length-1 ? a[tid+1] == 0 : a[0] == 0;
11          ensures tid != a.length-1 ? a[tid+1] == N*tid : a[0] == N*tid; */
12          {
13              /*@ loop_inv k >= 0 && k <= N;
14              loop_inv tid != a.length-1 ? Perm(a[tid+1], 1) : Perm(a[0], 1);
15              loop_inv tid != a.length-1 ? a[tid+1] == k*tid : a[0] == k*tid; */
16              for(int k = 0; k < N; k++) {
17                  if (tid != a.length-1) { a[tid+1] = a[tid+1] + tid; }
18                  else { a[0] = a[0] + tid; }
19              } } }
```

ALPINIST an Annotation-Aware GPU Program Optimizer



GPU Optimizations

GPU Optimizations

- Applicability of an optimization
- Complex decisions

Loop unrolling

```
1 void Host(int[] arr, int N){  
2     par kernel(tid=0..arr.length){  
3         int i = 0;  
4         while (i < N){  
5             int newInt = i;  
6             arr[tid] = arr[tid] + newInt;  
7             i = i + 1;  
8         }  
9     }  
10 }
```

```
1 void Host(int[] arr, int N){  
2     par kernel(tid=0..arr.length){  
3         int i = 0;  
4         int newInt = i;  
5         arr[tid] = arr[tid] + newInt;  
6         i = i + 1;  
7  
8         newInt = i;  
9         arr[tid] = arr[tid] + newInt;  
10        i = i + 1;  
11  
12        while (i < N){  
13            newInt = i;  
14            arr[tid] = arr[tid] + newInt;  
15            i = i + 1;  
16        }  
17    }  
18}
```

Loop unrolling with annotations

```
1  /*@ context N > 1; @*/
2  void Host(int[] arr, int N){
3      par kernel(tid=0..arr.length){
4          int i = 0;
5          /*@ loop_inv i >= 0 && i <= N;
6              loop_inv N > 1;
7              loop_inv Inv(i); @*/
8          while (i < N){
9              int newInt = i;
10             arr[tid] = arr[tid] + newInt;
11             i = i + 1;
12         }
13     }
14 }
```

```
1  /*@ context N > 1; @*/
2  void Host(int[] arr, int N){
3      par kernel(tid=0..a.length){
4          int i = 0;
5          int newInt = i;
6          arr[tid] = arr[tid] + newInt;
7          i = i + 1;
8          //@ assert i >= 1 && i <= N;
9          //@ assert N > 1;
10         //@ assert Inv(i);
11         newInt = i;
12         arr[tid] = arr[tid] + newInt;
13         i = i + 1;
14         /*@ loop_inv i >= 2 && i <= N;
15             loop_inv N > 1;
16             loop_inv Inv(i); @*/
17         while (i < N){
18             newInt = i;
19             arr[tid] = arr[tid] + newInt;
20             i = i + 1;
21         }
22     }
23 }
```

Kernel fusion

```
1 void Host(...){  
2     par kernel1(tid1 = 0..T) {  
3         a[tid1] = 2*b[tid1];  
4     }  
5  
6     par kernel2(tid2 = 0..T) {  
7         b[tid2] = a[tid2]+1;  
8     }  
9 }
```

```
1 void Host(...){  
2     par fused_kernels(tid = 0..T) {  
3         a[tid] = 2*b[tid];  
4         b[tid] = a[tid]+1;  
5     }  
6 }
```

Kernel fusion with annotations

```
1 void Host(...){  
2     par kernel1(tid1 = 0..T)  
3     /*@ context Perm(a[tid1], 1);  
4     context Perm(b[tid1], 1\2);@*/  
5     {  
6         a[tid1] = 2*b[tid1];  
7     }  
8  
9     par kernel2(tid2 = 0..T)  
10    /*@ context Perm(a[tid2], 1\2);  
11    context Perm(b[tid2], 1);@*/  
12    {  
13        b[tid2] = a[tid2]+1;  
14    }  
15}
```

Kernel fusion with annotations

```
1 void Host(...){  
2     par kernel1(tid1 = 0..T)  
3         /*@ context Perm(a[tid1], 1);  
4             context Perm(b[tid1], 1\2);@*/  
5     {  
6         a[tid1] = 2*b[tid1];  
7     }  
8  
9     par kernel2(tid2 = 0..T)  
10        /*@ context Perm(a[tid2], 1\2);  
11            context Perm(b[tid2], 1);@*/  
12        {  
13            b[tid2] = a[tid2]+1;  
14        }  
15    }
```

```
1 void Host(...){  
2     par fused_kernels(tid = 0..T)  
3     /*@ requires Perm(a[tid], 1);  
4     requires Perm(b[tid], 1\2);  
5     requires Perm(b[tid], 1\2);@*/  
6     {  
7         a[tid] = 2*b[tid];  
8         b[tid] = a[tid]+1;  
9     }  
10 }
```

```
1 ...
2 void simple_example(int[] a, int size, int N) {
3     par kernel1 (int tid = 0 .. a.length)
4         /*@ context Perm(a[tid], 1);
5         ensures a[tid] == 0; */
6     { a[tid] = 0; }
7
8     par kernel2 (int tid = 0 .. a.length)
9         /*@ context tid != a.length-1 ? Perm(a[tid+1], 1) : Perm(a[0], 1);
10        requires tid != a.length-1 ? a[tid+1] == 0 : a[0] == 0;
11        ensures tid != a.length-1 ? a[tid+1] == N*tid : a[0] == N*tid; */
12    {
13        /*@ loop_inv k >= 0 && k <= N;
14        loop_inv tid != a.length-1 ? Perm(a[tid+1], 1) : Perm(a[0], 1);
15        loop_inv tid != a.length-1 ? a[tid+1] == k*tid : a[0] == k*tid; */
16        for(int k = 0; k < N; k++) {
17            if (tid != a.length-1) { a[tid+1] = a[tid+1] + tid; }
18            else { a[0] = a[0] + tid; }
19        } } }
```

```
1 ...
2 void simple_example(int[] a, int size, int N) {
3     par kernel1 (int tid = 0 .. a.length)
4         /*@ context Perm(a[tid], 1);
5         ensures a[tid] == 0; */
6     { a[tid] = 0; }
7
8     par kernel2 (int tid = 0 .. a.length)
9         /*@ context tid != a.length-1 ? Perm(a[tid+1], 1) : Perm(a[0], 1);
10        requires tid != a.length-1 ? a[tid+1] == 0 : a[0] == 0;
11        ensures tid != a.length-1 ? a[tid+1] == N*tid : a[0] == N*tid; */
12    {
13        /*@ loop_inv k >= 0 && k <= N;
14        loop_inv tid != a.length-1 ? Perm(a[tid+1], 1) : Perm(a[0], 1);
15        loop_inv tid != a.length-1 ? a[tid+1] == k*tid : a[0] == k*tid; */
16        for(int k = 0; k < N; k++) {
17            if (tid != a.length-1) { a[tid+1] = a[tid+1] + tid; }
18            else { a[0] = a[0] + tid; }
19        } } }
```

Decision points in the perm-related precondition procedure

- ① Patterns are the same (i.e. thread to location mapping the same)

Decision points in the perm-related precondition procedure

- ① Patterns are the same (i.e. thread to location mapping the same)
- ② Permission for a location ≤ 1

Decision points in the perm-related precondition procedure

- ① Patterns are the same (i.e. thread to location mapping the same)
- ② Permission for a location ≤ 1
- ③ Only read permissions

Decision points in the perm-related precondition procedure

- ① Patterns are the same (i.e. thread to location mapping the same)
- ② Permission for a location ≤ 1
- ③ Only read permissions
- ④ One of the kernels has write permissions

Decision points in the perm-related precondition procedure

- ① Patterns are the same (i.e. thread to location mapping the same)
- ② Permission for a location ≤ 1
- ③ Only read permissions
- ④ One of the kernels has write permissions

Decision points in the perm-related precondition procedure

- ① Patterns are the same (i.e. thread to location mapping the same)
- ② Permission for a location ≤ 1
- ③ Only read permissions
- ④ One of the kernels has write permissions, data dependency


Supported GPU Optimizations in ALPINIST

- Loop unrolling
- Kernel fusion
- Tiling
- Iteration merging
- Matrix linearization
- Data prefetching

Supported GPU Optimizations in ALPINIST

- Loop unrolling
- Kernel fusion
- Tiling
- Iteration merging
- Matrix linearization
- Data prefetching

(Internal) Design of the Tool



Evaluation

Evaluation

- **Q1** test whether ALPINIST works on GPU programs.

Evaluation

- **Q1** test whether ALPINIST works on GPU programs.
- **Q2** investigate how long it takes for ALPINIST to transform GPU programs and how this affects the verification time.

Evaluation

- **Q1** test whether ALPINIST works on GPU programs.
- **Q2** investigate how long it takes for ALPINIST to transform GPU programs and how this affects the verification time.
- **Q3** investigate the usability of ALPINIST on real-world complex examples.

Q1 & Q3 Evaluation Suite

Q1 & Q3 Evaluation Suite

① Hand-made examples

Q1 & Q3 Evaluation Suite

- ① Hand-made examples
- ② Verified examples in VerCors

Q1 & Q3 Evaluation Suite

- ① Hand-made examples
- ② Verified examples in VerCors
- ③ Complex case studies in VerCors:
 - Two parallel prefix sum algorithms
 - Parallel stream compaction
 - Summed-area table algorithms
 - A variety of sorting algorithms
 - A solution to VerifyThis 2019 challenge 1
 - A Tic-Tac-Toe example

Q1 & Q3 Evaluation Suite

- ① Hand-made examples
- ② Verified examples in VerCors
- ③ Complex case studies in VerCors:
 - Two parallel prefix sum algorithms
 - Parallel stream compaction
 - Summed-area table algorithms
 - A variety of sorting algorithms
 - A solution to VerifyThis 2019 challenge 1
 - A Tic-Tac-Toe example

70 experiments in total across all 6 optimizations

Q2 Effect on verification time

Optimizations	Effect on Code/Annotations	Avg. Verification Time
---------------	-------------------------------	---------------------------

Q2 Effect on verification time

Optimizations	Effect on Code/Annotations	Avg. Verification Time
Loop unrolling	Generates code/annotations	+14%
Tiling	Generates code/annotations	+32%
Iteration merging	Generates code/annotations	+18%

Q2 Effect on verification time

Optimizations	Effect on Code/Annotations	Avg. Verification Time
Loop unrolling	Generates code/annotations	+14%
Tiling	Generates code/annotations	+32%
Iteration merging	Generates code/annotations	+18%
Kernel Fusion	Removes code/annotations	-23%

Q2 Effect on verification time

Optimizations	Effect on Code/Annotations	Avg. Verification Time
Loop unrolling	Generates code/annotations	+14%
Tiling	Generates code/annotations	+32%
Iteration merging	Generates code/annotations	+18%
Kernel Fusion	Removes code/annotations	-23%
Matrix Linearization	Replaces code/annotations	+1%
Data prefetching	Replaces code/annotations	-4%

To conclude

- ALPINIST: an Annotation-Aware GPU Program Optimizer

To conclude

- ALPINIST: an Annotation-Aware GPU Program Optimizer
- Six optimizations supported in the tool

To conclude

- ALPINIST: an Annotation-Aware GPU Program Optimizer
- Six optimizations supported in the tool
- Future work

To conclude

- ALPINIST: an Annotation-Aware GPU Program Optimizer
- Six optimizations supported in the tool
- Future work
 - More GPU optimizations

To conclude

- ALPINIST: an Annotation-Aware GPU Program Optimizer
- Six optimizations supported in the tool
- Future work
 - More GPU optimizations
 - Support for OpenCL and CUDA

ALPINIST an Annotation-Aware GPU Program Optimizer

